

## APPENDIX J

### REPRODUCIBILITY BUNDLE

Experiments ran on Windows 10 (64-bit, build 10.0.26100) with Python 3.11 on x64 CPUs. Planning proceeded for up to five iterations with a 300 s timeout and default temperature 0.7 (seeds not fixed). Measured end-to-end latency ranged from 30–300 s depending on itinerary complexity, with typical sessions involving 1–3 LLM calls per agent, 5–15 database queries, and 1.5k–7k tokens (0.5k–2k input; 1k–5k output), at an estimated cost of USD \$0.00015–\$0.0007 (GPT-4o). While introducing communication/memory overhead compared to single-agent baselines, the multi-agent design enabled modular specialization, parallelism, and feedback-driven refinement.